

Cuisine Bias in LLM Health Ratings Depends on How Culture Is Signalled

Anonymous ACL submission

Abstract

A cultural-bias audit must choose how to signal the cultural category. We show that this choice, not the model, determines how much bias the audit detects. Holding a dish’s ingredient list *byte-identical* and varying only the signal, we collect healthiness ratings from 20 open-weight instruction models across 11 families ($\approx 260,000$ generations). Naming the cuisine explicitly yields weak cross-model agreement on which cuisines are downrated ($\rho = +0.089$); conveying it implicitly through a dish name yields four times as much ($\rho = +0.365$), and the contrast is significant when tested directly ($\Delta\rho = +0.276$, $p = 0.0125$). The natural objection — that a dish name supplies a specific dish with its own nutrition prior, so models agree about dishes rather than cuisines — is refuted by a control using *invented* names that are morphologically diagnostic of a cuisine but denote nothing: agreement is undiminished ($\rho = +0.350$; $p = 0.441$ against real names). Two further ablations locate the effect precisely. A 0–100 scale reproduces it exactly, so it is not scale coarseness. But moving the same explicit signal from *before* to *after* the ingredient list raises its agreement more than fivefold ($+0.089 \rightarrow +0.493$) while preserving the cuisine ordering ($r = +0.87$) — detectable bias is a property of where the signal sits as much as what it says. We report reliability, congruence, power and magnitude analyses throughout, and are explicit that the underlying effects, though reproducible, are small. A single-channel audit measures how much bias its own elicitation can see.

1 Introduction

A standard bias audit asks a model to judge items differing only in a cultural marker and reports which group fares worse. The counterfactual looks clean. But the marker must be *delivered*, and several deliveries are equally natural: name the nationality, name the country, or simply name the

dish. We show these are not interchangeable. Two audits of the same models, both pre-registered and both internally consistent, differ fourfold in how much bias they detect — and a third variant, differing only in *where* the signal sits in the prompt, differs fivefold again.

We study food. Dietary advice is a high-volume consumer use of these models, and a systematic penalty against particular cuisines would be an allocational harm to populations already underserved by nutrition guidance. Food also isolates the methodological question unusually well: cultural identity can be conveyed several ways in ordinary language while the thing judged stays completely fixed. Prior work compares real dishes, which confounds cultural signal with genuine nutritional difference (Couscous, 2026); we remove that confound by construction.

Contributions.

1. A design in which nutrition is constant by *identity* rather than by matching, crossed over ingredient provenance so cuisine effects separate from ingredient–signal congruence (§3).
2. Implicit signalling elicits $\approx 4\times$ the cross-model agreement of explicit signalling, with the contrast tested *directly* rather than inferred from two separate within-channel tests (§4).
3. A pseudo-name control removing the referent dish and its nutrition prior; the effect is intact, ruling out shared dish knowledge (§5).
4. Five ablations — rating scale, signal position, instruction paraphrase, dish set, and forced choice — that locate the effect and expose two artefacts that would each have produced a spurious null (§6).
5. Reliability, congruence, magnitude and power analyses, including the finding

that a 13-model version of this study was underpowered to detect its own effect (§7).

2 Related work

Judge bias is documented for position, verbosity and self-preference (Silent Judge, 2025; Judge Bias, 2025), and writing style shifts automated assessment (Grading Bias, 2026). Cultural food knowledge has been probed by comparing real dishes (Couscous, 2026). Multilingual judging work reports lower-resource categories scored *more* generously (Lower-Resource, 2026), the direction we also find under explicit labels.

Two threads should be separated. The first concerns *what* judges are biased about. The second, much smaller, concerns *how bias is measured*, and it is here our result sits. Format sensitivity is documented in closely related forms: serialisation format changes tool-calling accuracy at fixed semantics (Notation Matters, 2026; Tool Calling, 2026); superficial variation in code changes how an LLM judges it (Code Cover, 2025); adding a schema-legal abstention option changes behaviour on items where abstention is unwarranted (Abstention Artifact, 2025; AbstentionBench, 2025; PhantomFill, 2026). In each the delivery mechanism, not the content, moves the measurement. Evaluation-validity work argues complementarily that reported effects can be artefacts of measurement design (Validity Audit, 2026; Great Pretender, 2026). Language and script carry costs invisible to content-only measurement (African Language Tax, 2026; Tokenizer Tax, 2026), and health communication is known to be style-sensitive (Health Comms, 2025). Recent judge studies evaluate many judges because single-model findings proved fragile; we argue the same fragility applies along a dimension currently unreported — the channel, and the position, through which the manipulated attribute is delivered.

3 Method

3.1 Notation and identification

Let $\mathcal{C}$ be the cuisines ($|\mathcal{C}|=6$), $\mathcal{I}$ the dish items ($|\mathcal{I}|=18$, three per cuisine) with provenance $\pi : \mathcal{I} \rightarrow \mathcal{C}$, and $\mathcal{M}$ the models ($|\mathcal{M}|=20$). A channel g maps a cuisine to a lead sentence $\ell_g(c)$; ℓ_0 is a neutral lead of matched form. Let $r_m(i, \ell, q, o)$ be the rating model m gives item i under lead ℓ , paraphrase $q \in [Q]$ and ingredient ordering $o \in [O]$,

$Q=O=3$. Decoding is greedy, so r_m is deterministic in its arguments.

With $\bar{r}_m(i, \ell) = \frac{1}{QO} \sum_{q,o} r_m(i, \ell, q, o)$, the channel effect is

$$\Delta_m^g(c) = \frac{1}{|\mathcal{I}_{\neq c}|} \sum_{i:\pi(i) \neq c} [\bar{r}_m(i, \ell_g(c)) - \bar{r}_m(i, \ell_0)]. \quad (1)$$

(1) identifies the cultural signal under one assumption, which the design enforces rather than assumes: $\ell_g(c)$ and ℓ_0 are evaluated on the *same* item, with a byte-identical ingredient list, in the same sentence slot. The food is therefore differentiated out exactly, and so is the mere presence of a lead sentence — a substantial effect in its own right, which we return to in §6. Restricting to $\pi(i) \neq c$ removes the congruent cell, which measures plausibility as much as culture; congruence is estimated separately, *per channel*, in §7.

3.2 The estimand

Collecting $E_{cm}^g = \Delta_m^g(c)$, the quantity of interest is whether models *agree* on the cuisine ordering,

$$\rho(E^g) = \binom{|\mathcal{M}|}{2}^{-1} \sum_{m < m'} \text{corr}(E_{\cdot m}^g, E_{\cdot m'}^g). \quad (2)$$

ρ is scale-free by construction: it is invariant to how strongly any individual model reacts, and asks only whether models rank cuisines alike. This is the right target for an audit question — a bias one model shows and no other is a property of that model, not of the class — but it means ρ must be read alongside effect magnitudes, which we report in §7.

3.3 Inference

Permutation. Significance of ρ comes from permuting cuisine labels *independently within each model column*, destroying any shared ordering while preserving each model’s marginal distribution of effects ($B=5000$). This assumes nothing about which ordering models would agree on.

Contrasts. A difference between a significant and a non-significant result is not a significant difference. Channel contrasts are therefore tested directly, by exchanging each model’s two channel vectors ($B=10^4$) and by a model-level bootstrap.

Bootstrap unit. Intervals resample the unit of generalisation: the item for effect sizes, the model for ρ and for mean ranks. Resampling individual

Table 1: Cross-model agreement on the cuisine ordering over 20 models and 11 families. Implicit signalling yields roughly four times the agreement of explicit signalling when the signal precedes the content.

Cultural-signal channel	ρ	perm. p
explicit demonym	+0.089	0.0166
explicit country	+0.050	0.0804
real dish name	+0.365	0.0000
pseudo dish name	+0.350	0.0000

ratings would treat $QO=9$ prompts about one dish as independent.

3.4 Materials

Each item is an ingredient list on a fixed template (grain, protein, vegetable, fat, aromatic, seasoning). Two rules are asserted programmatically: no ingredient string contains a country word or demonym, and all sets are vegetarian and alcohol-free, removing religious-taboo confounds by construction rather than post-hoc exclusion. Four channels share one sentence slot: *demonym*, *country*, *real dish name* (five per cuisine), and *pseudo dish name* (§5).

3.5 Models, admission and hygiene

20 open-weight instruction models across 11 families. A model is admitted only if $>80\%$ of its generations parse in *every* channel (Appendix A). This is not a formality. An earlier version capped generation at 8 new tokens, which truncated verbose models before they reached a digit: one model emitted “*I’d rate this meal as an*” and parsed at 0%, yet was still counted in the model total, contributing an all-NAN column that the correlation filter silently dropped. Raising the cap to 24 tokens recovers the same model at “...as an 8.” **Truncation length is a selection criterion on the model set, not an implementation detail.**

Noise floor. Since decoding is greedy, rating variation comes from the nuisance factors. We measure $\sigma = \text{median}_{m,i,\ell} \text{sd}_{q,o} r_m(i, \ell, q, o) = 0.500$ rating points over 28,795 cells. It is identical to two decimals in all four channels, so it is a property of the elicitation rather than of any channel.

4 Results

Both implicit channels produce strong agreement; both explicit channels produce weak agreement. The explicit demonym is *not* a clean null — at 20




Figure 1: Cross-model agreement by channel. Implicit signalling elicits roughly four times the agreement of explicit signalling, and invented names match real ones.

Table 2: The contrast tested directly, not inferred from two within-channel p -values.

Contrast	$\Delta\rho$	p	95% CI
real – demonym	+0.276	0.0125	[+0.03, +0.46]
pseudo – demonym	+0.261	0.0006	[+0.05, +0.42]
real – pseudo	+0.015	0.4407	[-0.18, +0.23]

models it reaches $p = 0.017$ — but it is four times smaller.

Tested directly (Table 2), real names exceed the demonym by +0.276 ($p = 0.0125$) and pseudo names by +0.261 ($p = 0.0006$); both intervals exclude zero. The two implicit channels do not differ from each other.

5 The pseudo-name control

A dish name supplies more than a cuisine: it supplies a dish with its own nutrition prior. Models agreeing that *Miso Shiru* is healthier than *Cas-soulet* would be shared factual knowledge about two dishes, not a shared ordering over cuisines. Leave-one-name-out cannot settle this, because the confound is shared by all five names within a cuisine.

We therefore built 30 *invented* names, five per cuisine, from each cuisine’s productive naming morphology (*Ruvellata al Forno*, *Habuki-jiru*, *Obe Ajiri*, *Sopa Zarquena*), so cuisine is recoverable from form while no referent dish exists and no nutrition prior can attach. A manipulation check confirms transmission: given only the invented name, models recover the intended cuisine at 0.759 mean accuracy against a chance level of 0.167 (real names: 0.927).

Agreement is undiminished ($\rho = +0.350$ vs.

+0.365; $p = 0.441$). Leave-one-name-out is also markedly more stable for invented names (median +0.998, min +0.953) than real ones (median +0.994, min +0.545) — exactly as expected if real names carry dish-specific information that invented ones do not, while the cuisine ordering does not depend on it.

6 Ablations

Scale is irrelevant; position is decisive. A 0–100 scale reproduces the main result almost exactly (+0.131 and +0.370), ruling out coarse quantisation as the reason explicit signalling looks weak. But moving the *same* explicit sentence from before the ingredient list to after it raises its agreement more than fivefold (+0.089 → +0.493), and the implicit channel rises too (+0.365 → +0.535). The cuisine ordering is preserved ($r = +0.87$ for the demonym), so this is the same bias measured more sensitively, not a different one. **How much cultural bias an audit detects depends on where the signal sits in the prompt.**

Not a wording or an ingredient set. Recomputing ρ within each instruction paraphrase and within each of the three dish sets leaves the channel gap intact in every slice (Table 4).

Forced choice is unusable, and would have produced a false null. Intended as a finer instrument — two dishes with identical ingredients, differing only in label — forced choice fails: 17 of 20 models are position-degenerate, picking one option almost regardless of content. Because presentation order is counterbalanced, those models’ win rates are forced to *exactly* 0.5 by construction. Including them would have yielded a flawless and entirely spurious null. Among the three usable models $\rho = +0.219$ and the ordering does not track the Likert one.

7 Robustness and magnitudes

Not an attenuation artefact. Implicit cells average five names, explicit cells one, so implicit effects could be less noisy and hence less attenuated. Two checks rule this out. Implicit effect vectors are more reliable, and the advantage survives correction (Table 6); and pinning the implicit channel to a *single* name — matching explicit precision exactly — still leaves it well ahead.

Congruence does not explain it. Measured per channel, congruence is no larger in the implicit




Figure 2: Magnitudes against the measured noise floor and the pre-registered smallest effect of interest.

channels, and nearly vanishes for pseudo-names, which have no canonical dish to clash with.

The effects are consistent but small. Cuisine-mean spreads are 0.12–0.33 rating points against a noise floor of 0.500, and only 7 of 120 model–cuisine cells exceed the pre-registered smallest effect of interest (Figure 2). We report a reproducible *ordering* of small effects and do not claim allocational harm follows at this size.

Power, and why n mattered. At 13 models the same contrast gave $\Delta\rho = +0.270$, $p = 0.081$, interval spanning zero. Table 7 shows why: at that size the design has power ≈ 0.53 . At 20 it reaches 0.76. An earlier version of this study could not have detected its own effect, and we report this because it is the difference between a claim and an impression.

8 Discussion

An audit signalling culture with a demonym placed before the content reports weak, barely-detectable cuisine bias. The same audit with the same sentence placed *after* the content reports five times more. An audit signalling culture through a dish name — real or invented — reports four times more again. All are pre-registered, internally consistent, and correct about their own elicitation. A single-channel audit measures a joint property of the model and the channel and reports it as a property of the model.

We therefore recommend that cultural-bias audits (i) hold substantive content identical rather than matched; (ii) report at least one explicit and one implicit channel; (iii) vary signal *position*, not only signal content; (iv) baseline against a neutral marker, not an absent one; (v) measure a phrasing

Table 3: **Rating scale is irrelevant; signal position is decisive.** A 0–100 scale reproduces the main result. Moving the same cultural signal to *after* the ingredient list raises explicit agreement more than fivefold, while preserving the cuisine ordering (rank r against the main design). Detectable bias is a property of where the signal sits, not only of what it says.

Elicitation variant	explicit demonym		real dish name	
	ρ	rank r	ρ	rank r
1–10, signal first	+0.089	+1.00	+0.365	+1.00
0–100 scale	+0.131	+0.93	+0.370	+1.00
1–10, signal <i>after</i> ingredients	+0.493	+0.87	+0.535	+0.98

Table 4: ρ recomputed within each instruction paraphrase and within each of the three dish sets. The channel gap holds in every slice; it is not carried by one wording or one set of ingredients.

Channel	by paraphrase			by dish set			pooled
	1	2	3	1	2	3	
explicit demonym	+0.072	+0.144	+0.099	+0.072	+0.180	+0.146	+0.089
real dish name	+0.353	+0.296	+0.506	+0.249	+0.368	+0.378	+0.365
pseudo dish name	+0.356	+0.291	+0.198	+0.382	+0.335	+0.166	+0.350

Table 5: Forced choice, intended as a finer instrument, is unusable here: position dominates content for most models, and counterbalancing forces their win rates to exactly 0.5.

Quantity	value
models run	20
position-degenerate	17
usable after screen	3
ρ on win rates	+0.219
rank r vs. Likert	-0.262

noise floor; (vi) verify the signal transmits; (vii) audit parse rates, since generation caps silently exclude verbose models; (viii) screen for position degeneracy before trusting forced choice; and (ix) use several model families, and enough of them — we show 13 was not enough to detect our own effect.

9 Limitations

Six cuisines, 18 items, 30 real and 30 invented names, one judgement dimension, one prompt family, open-weight models only. Each ρ averages pairwise correlations over 6-dimensional vectors. The effects are small — consistent orderings of largely sub-SESOI magnitudes — and we do not claim harm follows at this size. Invented names transmit cuisine less strongly than real ones (0.759 vs. 0.927), which may itself depress the pseudo-channel estimate; the control is conservative with respect to our own claim but not free of it. The manipulation check is model-based, not

human-normed: it establishes the signal is legible to models — the relevant quantity for a claim about model behaviour — but not that it carries the same associations for people. Five of 25 candidate models were excluded on parse rate, and we cannot rule out that verbose models differ systematically. Cuisines differ in cooking fat by design, so provenance-level comparisons stay confounded with real composition; only the channel effect holds ingredients fixed. Finally the healthiness construct is thin: a 1–10 rating may track richness, unfamiliarity or perceived processing. Ingredients are fixed, so whatever it tracks is applied to identical food, but we cannot say which construct is reported.

Ethics statement

This study measures model behaviour, not food. Nutrition is constant by construction, so **the design is silent on whether any cuisine is more or less healthy than any other**, and no result may be read as such a claim. We publish no ranking of cuisines; per-cuisine numbers are model bias magnitudes. National cuisine labels erase within-country variation, and we use specific named cuisines rather than a Western/non-Western binary. Ingredient sets are vegetarian and alcohol-free so religious-taboo pairings never arise. No human subjects were involved. Model outputs are measurements, not dietary advice. Code, prompts, stimuli, the exclusion list and all raw ratings are released.

Table 6: **Reliability**: the implicit advantage survives correction, so it is not an attenuation artefact. **Congruence**: no larger in implicit channels, and nearly absent for pseudo-names, which have no canonical dish to clash with. **Magnitude**: consistent orderings of small effects.

Channel	Reliability			Congruence		Magnitude		
	inter-item r	S-B	ρ corr.	effect	95% CI	spread	max	>SESOI
explicit demonym	+0.45	+0.71	+0.125	+0.285	[+0.19, +0.40]	0.183	1.66	4/120
explicit country	+0.42	+0.68	+0.074	+0.263	[+0.17, +0.39]	0.120	1.74	4/120
real dish name	+0.70	+0.88	+0.416	+0.244	[+0.16, +0.33]	0.327	1.67	7/120
pseudo dish name	+0.74	+0.90	+0.391	+0.077	[+0.05, +0.11]	0.178	0.82	0/120

Table 7: Power for the name-vs-demonym contrast, by bootstrap over models. An earlier 13-model version of this study could not have detected its own effect.

Models n	power to detect $\Delta\rho$
8	0.35
12	0.53
16	0.61
20	0.76
24	0.86
30	0.91

References

The Silent Judge: Unacknowledged Shortcut Bias in LLM-as-a-Judge. arXiv:2509.26072, 2025.

Evaluating and Mitigating LLM-as-a-judge Bias in Communication Systems. arXiv:2510.12462, 2025.

Implicit Grading Bias in Large Language Models. arXiv:2603.18765, 2026.

Can LLMs Cook Jamaican Couscous? arXiv:2602.10964, 2026.

Lower-Resource, Higher Scores. arXiv:2607.14480, 2026.

LLM Abstinence Can Be a Prompt Artifact. arXiv:2507.16199, 2025.

AbstinenceBench. arXiv:2506.09038, 2025.

PhantomFill: Schema-Constrained Fabrication. arXiv:2607.20492, 2026.

Notation Matters. arXiv:2605.29676, 2026.

The Bitter Lesson of Tool Calling. arXiv:2608.06370, 2026.

Don't Judge Code by Its Cover. arXiv:2505.16222, 2025.

The Great Pretender. arXiv:2605.14418, 2026.

Benchmarking the Benchmarks. arXiv:2607.02577, 2026.

Communication Styles and Reader Preferences in Explaining Health Information. arXiv:2505.08143, 2025.

The African Language Tax. arXiv:2606.24460, 2026.

The Tokenizer Tax. arXiv:2607.24276, 2026.

Qwen Team. *Qwen2.5 Technical Report*. arXiv:2412.15115, 2024.

Grattafiori et al. *The Llama 3 Herd of Models*. arXiv:2407.21783, 2024.

Gemma Team. *Gemma 2*. arXiv:2408.00118, 2024.

Jiang et al. *Mistral 7B*. arXiv:2310.06825, 2023.

OLMo Team. *2 OLMo 2 Furious*. arXiv:2501.00656, 2025.

Abdin et al. *Phi-3 Technical Report*. arXiv:2404.14219, 2024.

DeepSeek-AI. *DeepSeek LLM*. arXiv:2401.02954, 2024.

SmolLM2: When Smol Goes Big. arXiv:2502.02737, 2025.

01.AI. *Yi: Open Foundation Models*. arXiv:2403.04652, 2024.

A Parse rates and admission

B Signal position




Figure 3: Signal position dominates rating scale.

C Power

D Per-model effects

Table 8: Fraction of generations yielding a parseable rating, per admitted model. Models below 0.80 in any channel were excluded. An earlier 8-token generation cap truncated verbose models before they reached a digit and silently removed them from the study.

Model	demonym	country	name	name
deepseek-llm-7b-chat	0.87	0.91	0.89	0.89
granite-3.0-8b-instruct	0.99	1.00	0.98	0.96
Meta-Llama-3.1-8B-Instruct	1.00	1.00	0.98	1.00
Yi-1.5-9B-Chat	1.00	1.00	1.00	1.00
SmolLM2-1.7B-Instruct	1.00	1.00	1.00	1.00
Hermes-3-Llama-3.1-8B	1.00	1.00	1.00	1.00
Qwen2.5-1.5B-Instruct	1.00	1.00	1.00	1.00
Qwen2.5-14B-Instruct	1.00	1.00	1.00	1.00
Qwen2.5-7B-Instruct	1.00	1.00	1.00	1.00
OLMo-2-1124-13B-Instruct	1.00	1.00	1.00	1.00
OLMo-2-1124-7B-Instruct	1.00	1.00	1.00	1.00
granite-3.1-8b-instruct	1.00	1.00	1.00	1.00
Phi-3.5-mini-instruct	1.00	1.00	1.00	1.00
Phi-4-mini-instruct	1.00	1.00	1.00	1.00
OpenHermes-2.5-Mistral-7B	1.00	1.00	1.00	1.00
Falcon3-10B-Instruct	1.00	1.00	1.00	1.00
Falcon3-7B-Instruct	1.00	1.00	1.00	1.00
Llama-3.2-3B-Instruct	1.00	1.00	1.00	1.00
gemma-2-2b-it	1.00	1.00	1.00	1.00
gemma-2-9b-it	1.00	1.00	1.00	1.00




Figure 4: Power for the name-vs-demonym contrast by number of models.

Table 9: Per-model cuisine effect against the neutral lead, explicit demonym channel, in rating points, all 20 admitted models. Individual models are not inert; what varies is the direction.

Model	Ital	Fren	Japa	Indi	Nige	Mexi	spread
Yi-1.5-9B-Chat	-0.16	-0.16	-0.30	-0.04	+0.10	-0.09	0.39
SmolLM2-1.7B-Instruct	+0.07	+0.13	+0.39	+0.15	-0.01	+0.21	0.40
Hermes-3-Llama-3.1-8B	+0.01	+0.04	+0.04	+0.01	+0.01	-0.00	0.04
Meta-Llama-3.1-8B-Instru	-0.26	-0.07	-0.08	-0.00	+0.01	-0.07	0.26
Qwen2.5-1.5B-Instruct	-0.36	-0.26	-0.09	+0.06	-0.08	-0.30	0.41
Qwen2.5-14B-Instruct	-1.56	-0.90	-1.66	-0.57	-0.41	-1.10	1.24
Qwen2.5-7B-Instruct	-1.01	-0.93	-0.20	-0.39	-0.01	-0.37	1.00
OLMo-2-1124-13B-Instruct	-0.21	-0.19	-0.19	-0.11	-0.07	-0.15	0.15
OLMo-2-1124-7B-Instruct	-0.01	-0.16	-0.07	-0.21	-0.21	-0.22	0.21
deepseek-llm-7b-chat	-0.77	-0.53	-0.46	-0.75	-0.64	-0.73	0.32
granite-3.0-8b-instruct	-0.16	-0.21	-0.22	-0.08	-0.16	-0.07	0.15
granite-3.1-8b-instruct	-0.46	-0.36	-0.91	-0.56	-0.30	-0.48	0.61
Phi-3.5-mini-instruct	-0.16	-0.16	-0.02	-0.05	-0.07	-0.16	0.13
Phi-4-mini-instruct	-0.27	-0.33	-0.39	-0.27	-0.01	-0.39	0.39
OpenHermes-2.5-Mistral-7	-0.27	-0.17	-0.33	-0.01	+0.01	-0.05	0.33
Falcon3-10B-Instruct	+0.08	-0.17	-0.19	-0.19	-0.27	+0.03	0.36
Falcon3-7B-Instruct	-0.06	-0.11	-0.27	-0.16	-0.07	-0.09	0.21
Llama-3.2-3B-Instruct	-0.27	-0.18	-0.21	-0.01	-0.21	-0.24	0.26
gemma-2-2b-it	-0.47	-0.17	-0.48	-0.18	-0.19	-0.52	0.35
gemma-2-9b-it	-0.35	-0.12	-0.59	-0.17	-0.39	-0.33	0.47

Table 10: Per-model cuisine effect against the neutral lead, real dish name channel, in rating points, all 20 admitted models. Individual models are not inert; what varies is the direction.

Model	Ital	Fren	Japa	Indi	Nige	Mexi	spread
Yi-1.5-9B-Chat	-0.44	-0.54	-0.29	-0.48	-0.13	-0.24	0.41
SmolLM2-1.7B-Instruct	-0.37	-0.37	-0.06	-0.23	-0.79	-0.35	0.73
Hermes-3-Llama-3.1-8B	-0.18	-0.14	-0.00	-0.12	-0.11	-0.11	0.17
Meta-Llama-3.1-8B-Instru	-0.16	-0.18	-0.11	-0.22	-0.29	-0.16	0.17
Qwen2.5-1.5B-Instruct	-0.63	-0.40	-0.24	-0.33	-0.28	-0.39	0.40
Qwen2.5-14B-Instruct	-1.42	-1.35	-0.92	-1.53	-0.27	-0.83	1.26
Qwen2.5-7B-Instruct	-0.49	-0.81	-0.05	-0.33	+0.03	-0.11	0.84
OLMo-2-1124-13B-Instruct	-0.25	-0.24	-0.21	-0.25	-0.22	-0.02	0.24
OLMo-2-1124-7B-Instruct	-0.55	-0.55	-0.19	-0.42	-0.37	-0.32	0.36
deepseek-llm-7b-chat	-0.96	-1.03	-0.55	-0.78	-0.62	-1.11	0.56
granite-3.0-8b-instruct	-0.72	-0.36	-0.50	-0.51	-0.29	-0.62	0.43
granite-3.1-8b-instruct	-0.81	-0.65	-0.35	-0.76	-0.48	-0.55	0.46
Phi-3.5-mini-instruct	-0.34	-0.69	+0.12	-0.45	-0.38	-0.46	0.81
Phi-4-mini-instruct	-1.53	-1.67	-0.13	-0.80	-0.19	-0.42	1.53
OpenHermes-2.5-Mistral-7	-0.25	-0.25	-0.01	-0.17	+0.01	-0.00	0.26
Falcon3-10B-Instruct	-0.19	-0.28	-0.24	-0.35	-0.27	-0.16	0.18
Falcon3-7B-Instruct	-0.28	-0.31	-0.33	-0.49	-0.22	-0.23	0.27
Llama-3.2-3B-Instruct	-0.56	-0.60	-0.28	-0.32	-0.22	-0.42	0.39
gemma-2-2b-it	-0.36	-0.50	-0.08	-0.09	+0.04	-0.02	0.54
gemma-2-9b-it	-0.72	-0.48	-0.46	-0.37	-0.33	-0.33	0.39